

QUANTAREON

Investor Memo — Seed Round, \$500K

Founder	Vlad M., sole founder, QUANTAREON Labs (Tbilisi, Georgia)
Round	\$500K seed, 18-month runway
Stage	Methodology published (DOI), one prototype case on GitHub, product in specification
Date	May 2026

One-liner

QUANTAREON is building the **ConsciousAI engine** — a deliberation layer above any LLM that reduces reactivity at the architectural level. One engine, three market outlets: certification, coding, methodology licensing.

The Problem

Modern LLMs share one structural flaw: they are reactive. They emit the statistically most probable next token, not the most correct one. This produces:

- **Hallucinations** — under output pressure, the model invents instead of pausing.
- **Scope drift** — tasks expand beyond what was actually requested.
- **Sycophancy** — models agree under social pressure rather than verifying.
- **Jailbreak surface** — the first impulse to an adversarial prompt is the wrong one.

Every existing alignment technique — RLHF, Constitutional AI, chain-of-thought, ReAct, Reflexion — operates **after** the model has committed to its first impulse. The trajectory is already locked.

The Solution

The ConsciousAI Protocol is a system-prompt layer that bends the generation trajectory **before** the first token. Two parallel layers:

Static Identity (WHO)	Dynamic Action (HOW)
VL-Framework: identity seal, values, depth rules, memory cycle	Conscious Cycle: per-request gating — Initiation, Reflection, Ada
R-Cycle: 5 deliberation states, R0 as default background process	4-Value Review: Proposal → Risk Check → Scope Guard →

Cross-model. Same behavioral signature on closed-weight (Western) and open-weight (Eastern) frontier models — the work happens at the system-prompt layer, not inside any one model's internals.

The Proof Point

quantarion-router — multi-provider LLM router with fallback and circuit breaker, written solo end-to-end under this protocol in a **single session**: specification, implementation, and the full test suite — **124 passing tests, mypy --strict**, zero required runtime dependencies. The repo is public;

commit history shows the trajectory.

The Market

One engine, three market outlets. Each outlet stands on its own; together they compound.

Outlet	Buyer	What we sell
QUANTARION Audit	Compliance officers, regulators, independent third-party evaluator for AI agents. Built around EU AI Act	AI safety tests
QUANTARION Coding	Developers shipping production code with coding plans	Router with the protocol applied through the development
ConsciousAI API	Platform integrators, B2B SaaS	Methodology licensing: protocol-as-a-layer applied to client's existing LLMs

Timing. EU AI Act enforcement begins Q3 2026. Articles 14-15 create regulatory demand for independent agent evaluators. QUANTARION Audit lands into that demand at audit.quantareon.com.

Why Us

- **Methodology already public.** CC BY-NC-ND 4.0, DOI-anchored. The protocol exists, runs in production, and is replicable.
- **Proof point already shipped.** quantarion-router (124 tests, single session) shows the productivity profile of the protocol.
- **Cross-model from day one.** No vendor lock-in. Sufficient instruction-following is the only constraint.
- **Engineering-first.** The founder is the conductor, not just a researcher. Specs, tests, types, ship.

Use of Funds — \$500K, 18-month runway

Category	% of round	What it buys
Founder + 1 senior engineer salary, 6 months	30%	Two-person team to close the audit MVP
Compute + cross-model benchmarking (HarmBench, AdvBench, TruthfulQA)	20%	Formalize the empirical claim
Audit platform engineering (audit.quantareon.com)	20%	Test category spec to live evaluator
Coding platform core engine	15%	Productize the router pattern at scale
Legal (Delaware C-Corp flip, IP)	5%	Capped at \$25K, deferred until closing
Buffer / runway extension	10%	12-month reserve to extend to 18 months

Traction

- **Methodology published** — CC BY-NC-ND 4.0, DOI 10.5281/zenodo.19858814.
- **quantarion-router** — live on GitHub, 124 passing tests, mypy --strict.
- **Manifund campaign** — live at manifund.org, goal \$150,700, closes June 12, 2026.
- **LTFF application** — submitted May 2, 2026, under evaluation.
- **Public engineering report** — published on EA Forum, May 2026.

Roadmap (18 months)

Quarter	Milestone
Q3 2026	Audit MVP live at audit.quantareon.com. First 3 pilot clients.
Q4 2026	Cross-model benchmarks published. Coding platform alpha.
Q1 2027	Audit revenue traction. Coding beta, first 10 paying teams.
Q2 2027	Coding platform GA. First Conscious AI API enterprise contract. Series A readiness.

Structure

- **QUANTAREON Labs** (Tbilisi, Georgia) — operations, IP, day-to-day work.
- **QUANTAREON Inc.** (Delaware C-Corp) — holding company, recipient of investment, equity issuer.
- Standard flip on closing. Legal fees deferred until closing, reimbursed from investment proceeds (capped at \$25K).
- Founder remains based in Tbilisi.

Ask

\$500,000 seed. SAFE or priced round, terms open to discussion. Single lead preferred. **Macroscopic Ventures** is our first-choice partner — your existing portfolio (Gray Swan, Halcyon, Mythos, RunRL) is the exact ecosystem this work belongs in.

Contact

Email	vlad@quantareon.com
Telegram	@quantareon
Platform	quantareon.com
Methodology	github.com/makx518-ui/consciousness-protocol (DOI 10.5281/zenodo.19858814)

“The first impulse is noise. The second is signal. Silence → Clarity → Action.”

— Vlad M., founder, QUANTAREON Labs